

NON-CLASSICAL SEARCH ALGORITHMS BY STUART RUSSELL

MODIFIED BY JACEK MALEC FOR LTH LECTURES
JANUARY 28TH, 2013

CHAPTER 4 OF AIMA

© Stuart Russell

Chapter 4 of AIMA 1

Outline

- ◇ Hill-climbing
- ◇ Simulated annealing (briefly)
- ◇ Genetic algorithms (briefly)
- ◇ Local search in continuous spaces (briefly)
- ◇ (NEW!) Searching with nondeterministic actions (briefly)
- ◇ (NEW!) Searching with partial observations (briefly)
- ◇ (NEW!) Online search and unknown environments (briefly)

© Stuart Russell

Chapter 4 of AIMA 2

Iterative improvement algorithms

In many optimization problems, **path** is irrelevant;
the goal state itself is the solution

Then state space = set of “complete” configurations;
find **optimal** configuration, e.g., TSP
or, find configuration satisfying constraints, e.g., timetable

In such cases, can use **iterative improvement** algorithms;
keep a single “current” state, try to improve it


Constant space, suitable for online as well as offline search

© Stuart Russell

Chapter 4 of AIMA 3

Example: Travelling Salesperson Problem

Start with any complete tour, perform pairwise exchanges



Variants of this approach get within 1% of optimal very quickly with thousands of cities


© Stuart Russell

Chapter 4 of AIMA 4

Example: n -queens

Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal

Move a queen to reduce number of conflicts



Almost always solves n -queens problems almost instantaneously
for very large n , e.g., $n = 1\text{million}$

© Stuart Russell

Chapter 4 of AIMA 5

Hill-climbing (or gradient ascent/descent)

“Like climbing Everest in thick fog with amnesia”


```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                  neighbor, a node
  current ← MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbor ← a highest-valued successor of current
    if VALUE[neighbor] ≤ VALUE[current] then return STATE[current]
    current ← neighbor
  end
```

© Stuart Russell

Chapter 4 of AIMA 6

Hill-climbing contd.

Useful to consider **state space landscape**



Random-restart hill climbing overcomes local maxima—trivially complete

Random sideways moves ☹️ escape from shoulders 🚫 loop on flat maxima

Simulated annealing

Idea: escape local maxima by allowing some ‘bad’ moves
but **gradually decrease their size and frequency**

```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  inputs: problem, a problem
         schedule, a mapping from time to “temperature”
  local variables: current, a node
                  next, a node
                  T, a “temperature” controlling prob. of downward steps

  current ← MAKE-NODE(INITIAL-STATE[problem])
  for t ← 1 to ∞ do
    T ← schedule[t]
    if T = 0 then return current
    next ← a randomly selected successor of current
    ΔE ← VALUE[next] − VALUE[current]
    if ΔE > 0 then current ← next
    else current ← next only with probability  $e^{\Delta E/T}$ 
```

Properties of simulated annealing

At fixed “temperature” T , state occupation probability reaches Boltzman distribution

$$p(x) = \alpha e^{\frac{E(x)}{T}}$$

T decreased slowly enough \implies always reach best state x^*

because $e^{\frac{E(x^*)}{T}} / e^{\frac{E(x)}{T}} = e^{\frac{E(x^*) - E(x)}{T}} \gg 1$ for small T

Is this necessarily an interesting guarantee??

Devised by Metropolis et al., 1953, for physical process modelling

Widely used in VLSI layout, airline scheduling, etc.

Local beam search

Idea: keep k states instead of 1; choose top k of all their successors

Not the same as k searches run in parallel!

Searches that find good states recruit other searches to join them


Problem: quite often, all k states end up on same local hill

Idea: choose k successors randomly, biased towards good ones

Observe the close analogy to natural selection!

Genetic algorithms

= stochastic local beam search + generate successors from **pairs** of states




Fitness Selection Pairs Cross-Over Mutation

Genetic algorithms contd.

GAs require states encoded as strings (GPs use programs)

Crossover helps iff substrings are meaningful components



GAs ≠ evolution: e.g., real genes encode replication machinery!

Continuous state spaces

Suppose we want to site three robot battery loading stations in the hospital:

- 6-D state space defined by $(x_1, y_1), (x_2, y_2), (x_3, y_3)$
- objective function $f(x_1, y_1, x_2, y_2, x_3, y_3) =$ sum of squared distances from each location to nearest loading station

Discretization methods turn continuous space into discrete space, e.g., **empirical gradient** considers $\pm\delta$ change in each coordinate

Gradient methods compute

$$\nabla f = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial y_1}, \frac{\partial f}{\partial x_2}, \frac{\partial f}{\partial y_2}, \frac{\partial f}{\partial x_3}, \frac{\partial f}{\partial y_3} \right)$$


to increase/reduce f , e.g., by $\mathbf{x} \leftarrow \mathbf{x} + \alpha \nabla f(\mathbf{x})$

Sometimes can solve for $\nabla f(\mathbf{x}) = 0$ exactly (e.g., with one location).

Newton-Raphson (1664, 1690) iterates $\mathbf{x} \leftarrow \mathbf{x} - \mathbf{H}_f^{-1}(\mathbf{x}) \nabla f(\mathbf{x})$

to solve $\nabla f(\mathbf{x}) = 0$, where $\mathbf{H}_{ij} = \partial^2 f / \partial x_i \partial x_j$

Searching with nondeterministic actions




Erratic vacuum world: modified **Suck**;

Slippery vacuum world: modified **Right** and **Left**.

Searching with nondeterministic actions

And-or search trees


For the erratic case:



Searching with nondeterministic actions

And-or search trees

For the slippery case:



Searching with nondeterministic actions

function AND-OR-GRAPH-SEARCH(*problem*) *returns a cond. plan, or failure*
 OR-SEARCH(*problem*.INITIAL-STATE,*problem*,[])

function OR-SEARCH(*state*,*problem*,*path*) *returns a conditional plan or failure*
 if *problem*.GOAL-TEST(*state*) then return the empty plan
 if *state* is on *path* then return failure
 for each *action* in *problem*.ACTIONS(*state*) do
 plan ← AND-SEARCH(RESULTS(*state*,*action*),*problem*,[*state* | *path*])
 if *plan* ≠ failure then return [*action* | *plan*]
 return failure


function AND-SEARCH(*states*,*problem*,*path*) *returns a conditional plan or failure*
 for each *s_i* in *states* do
 plan_i ← OR-SEARCH(*s_i*,*problem*,*path*)
 if *plan_i* = failure then return failure
 return [if *s₁* then *plan₁* else if *s₂* then *plan₂* else if ... *plan_{n-1}* else *plan_n*]

Searching with partial observations

- ◇ no-information case:
 - sensorless problem, or
 - conformant problem
- ◇ state-space search is made in **belief space**
- ◇ Problem solving: and-or search!

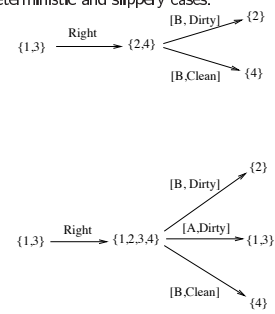
Searching with partial observations

Deterministic case:




Searching with partial observations

Local sensing, deterministic and slippery cases:



Searching with partial observations

Planning for the local sensing case:



Online search and unknown environments

Interleaving computations and actions:

- ◇ act
- ◇ observe the results
- ◇ find out (compute) next action

Useful in dynamic domains.

Online search usually exploits locality of depth-first-like methods.

- ◇ random walk
- ◇ modified hill-climbing
- ◇ Learning Real-Time A* (LRTA*)

optimism under uncertainty
(unexplored areas assumed to lead to goal with least possible cost)